

01 Problem addressed and use case reformulation

UC2 - BDR Outreach & Campaign Personalization for Mirakl Connect

Today, Mirakl sales teams spend too much time preparing prospecting instead of doing what they do best: converting.

Seller sourcing is still too manual: web search, catalog review, marketplace presence verification.

Seller -> marketplace matching is not industrialized nor explainable enough to be replayed at scale.

Email personalization is slow: either we keep quality and lose volume, or we increase volume and degrade relevance.

Business impact: slower seller pipeline, incomplete market coverage, and BDR time consumed by repetitive tasks.

Simple use case reformulation

Build an agentic system that detects the right brands, scores them against the right marketplaces, generates personalized outreach, and lets the human validate before sending.

Our answer: an end-to-end engine that turns scattered public signals into commercial opportunities ready to be handled.

147

scored brands

7

marketplaces

P1

prioritization

1. Discover

Find relevant sellers among hundreds of public brands.

2. Decide

Quickly understand which marketplace is the best fit and why.

3. Contact

Produce a credible, personalized email ready to send.

Our stance: the agent prepares the volume, the human keeps the commercial judgment

Core stance

The BDR should no longer do the research.

Their real value is in the conversation, qualification, negotiation, and closing.
AI handles the preparation, not the relationship.

1. Human-in-the-loop

No email is sent without human validation. We automate the preparation, not the responsibility.

2. Auditable matching

Each score must be explained by readable signals: category, geography, price tier, readiness, marketplace presence.

Why it matters for Mirakl

The system stays credible for the sales teams.

The matching can be defended in front of operators.

The prototype is already designed for real use, not just for a demo.

3. Segment before personalizing


We first adapt the tone, angle, and arguments to the brand archetype, then personalize in detail.

4. Cross-reference multiple signals

Quality comes from cross-referencing catalog + marketplace fit + positioning + public presence, not from a single isolated signal.

03 Synthetic technical architecture (macro view)

A simple chain: capture -> enrich -> score -> write -> validate -> steer



Stack used in the prototype: n8n for orchestration, Dust / LLM agents for reasoning and writing, TypeScript app for the business interface.

04

Annotated walkthrough - step 1: prioritize brands

Home view: the BDR starts with a shortlist instead of a blank page

What this screen shows

Portfolio overview: 147 brands, 7 marketplaces, and campaign targets.
Table of priority brands with score and P1 priority level.
Simple entry point to launch a campaign without manually scanning the entire market.

MIRAKL CONNECT

Marketplace Growth Engine

NAVIGATION

Accueil

Boîte de réception

Marques

Campagnes

Paramètres

localhost

Lancer la campagne

147

MARQUES TOTALES

7

MARKETPLACES

147

CIBLES CAMPAGNE

0

EMAILS ENVOYÉS

0

RÉUNIONS OBTENUES

0%

TAUX DE RÉPONSE

MARQUES PRIORITAIRES

BRAND	MARKETPLACE	SCORE	PRIORITY
Mothers Day (Agua by Agua Bendita)	Nordstrom	93.7	P1 — Fort potentiel
Maje	Galleries Lafayette	92.0	P1 — Immédiat
Okaidi	La Redoute	92.0	P1 — Immédiat
Sandro	Nordstrom	92.0	P1 — Immédiat
Doen	Debenhams	91.0	P1 — Fort potentiel
Aigle	John Lewis	91.0	P1 — Immédiat
Lacoste	Nordstrom	91.0	P1 — Immédiat
Petit Bateau	John Lewis	91.0	P1 — Immédiat

Takeaway

The product promise is visible immediately: less searching, more deciding.

05

Annotated walkthrough - step 2: explain the matching

Brand view: understand which marketplace is best and why

What the BDR understands here

The score is not opaque: it is broken down by business criteria.

Each marketplace is compared with a clear score and priority.

The salesperson can justify their target choice even before writing the first message.

MIRAKL CONNECT

Marketplace Growth Engine

NAVIGATION

Accueil

Boîte de réception

Marques

Campagnes

Paramètres

← Intelligence Marques

Okaidi

France · Kids value <https://www.okaidi.fr/>

P1 — Immédiat

PROFIL DE STRATÉGIE

Balanced (default)

92.0

La Redoute

P1

89.0

Debenhams

P1

Détail du score

Recommandations

Aperçu email

Profil de strategie

Fit catégorie

30

Fit géographique

12

Échelle commerciale

15

Readiness opérationnelle

13

Fit positionnement

12

Incrémentalité

8

Story durabilité

5

MARKETPLACE	SCORE	PRIORITY
La Redoute	92.0	P1
Debenhams	89.0	P1
John Lewis	85.0	P2
Nordstrom	83.0	P2
Bloomingdale's	80.0	P3
Galleries Lafayette	79.0	P3

Takeaway

Matching becomes defensible, reusable, and much faster to review.


Launch modal: brand selection and strategic profile choice

What this step enables

Pick a strategy profile to adapt the outreach angle.

Select the brands to contact and name the campaign.

Prepare a coherent batch of sendouts without any technical handling.



07

Annotated walkthrough - step 4: steer the execution


Campaign dashboard: track outreach volumes, statuses, and KPIs

What the dashboard measures

How many targets were loaded, contacted, and replied.

Status per brand: queued, paused, stopped, progress through the campaign.

A control cockpit to track the agent's performance.



Takeaway

The system is not just generative: it is also operable and measurable.


Inbox: final human control before any output

Why this step matters

All previews land in a centralized validation area.

The human can review, adjust, accept or reject before sending.

Reputational risk stays under control even at high volume.



MIRAKL CONNECT
Marketplace Growth Engine

NAVIGATION

- Accueil
- Boîte de réception**
- Marques
- Campagnes
- Paramètres

VALIDATION DES ENVOIS

23 avril 2026

Boîte de réception

Relisez les aperçus préparés par n8n, ajustez le ton si besoin, puis approuvez l'envoi. Rien ne part sans votre validation.

I • À VALIDER	II • MODIFIÉS	III • ENVOYÉS	IV • REFUSÉS	V • ÉCHECS
00	00	00	00	00

Note — Vous pouvez relire, ajuster le texte si besoin, puis valider l'envoi. Aucun email ne part tant que vous n'avez pas cliqué sur « Envoyer ».

Rechercher (marque, email, objet)...

00 APERÇU

Tirage en attente


Aucun aperçu en attente. n8n n'a encore rien envoyé à valider.

Takeaway

Trust does not come from autonomous AI, but from supervised AI.

09 What the agent does, what the human keeps

Human-in-the-loop by design: augmented capacity, responsibility preserved



Current limitations and identified improvement areas

Honest position: what already works, what must be strengthened to move from prototype to run

Uneven signal coverage

Some verticals or small brands have less usable public data.

Add fresh signals

Funding, hiring, catalog launches, launch signals.

Low real-time intent

The prototype mainly reads static signals; it still captures few hot events.

Closed CRM loop

Feed edits, replies, wins, and losses back into the scoring.

Multilingual quality to improve

Personalization remains stronger in the best-covered markets and tones.

Systematic A/B testing

Measure hook, subject, CTA, and angle per segment.

Human validation = bottleneck

Essential today, but it caps throughput at scale.

Levels of autonomy

Move from full validation to confidence-threshold validation.

The prototype is therefore credible to demonstrate value. Scaling will mainly depend on signal quality, the feedback loop, and the accepted level of automation.

11

Estimated run cost of the solution (FinOps)

MVP order of magnitude - indicative estimate to understand the variable cost of a campaign

Reading assumption

Example of a 100-prospect campaign processed with light enrichment, scoring, email generation, and human validation. Figures are indicative and must be recalibrated in production.

Component	Role	Estimated cost / 100 prospects
Collection & enrichment	Public scraping, cleaning, signal consolidation	€2 - €4
Reasoning & scoring	Agents / LLM for segmentation and explainable matching	€2 - €3
Email generation	Main email + variant + parsing	€2 - €4
Orchestration & app	n8n, light storage, TypeScript interface	€1 - €2

€7 - €13

variable cost / 100 prospects

€0.07 - €0.13

cost / prospect

Very low

marginal cost per email

Business takeaway

Even with human validation, machine cost stays low compared to the BDR time saved. The real FinOps topic will be the trade-off between quality / volume / level of autonomy.

What Mirakl can concretely do with this

Direct value: turn a hackathon prototype into an operational commercial capability

1. Cover more of the market

Move from very hand-crafted prospecting to much broader seller coverage without growing the team at the same pace.

2. Standardize quality

Make matching and outreach more consistent across BDRs, markets, and campaigns.

3. Reduce unproductive time

Give sales reps back time for calls, objections, and closing.

4. Capitalize on signals

Turn public data and human edits into a commercial learning engine.

Realistic trajectory after the hackathon

Pilot with a few operators, CRM connection, signal enrichment, learning from validations, and gradual ramp-up of autonomy.

In one sentence: Mirakl can turn this prototype into an AI-assisted seller growth engine, controlled by the teams and measurable over time.